

PAIHENG XU

Email: paiheng@umd.edu ◊ Homepage: paihengxu.github.io ◊ Telephone: (410) · 900 · 3657 ◊ Github: [paihengxu](https://github.com/paihengxu)

RESEARCH INTEREST

Natural Language Processing, Computational Social Science, Causal Inference

EDUCATION

University of Maryland, College Park, USA

2021 - present

Ph.D. Student in Computer Science

Johns Hopkins University, Baltimore, USA

2018 - 2020

M.S.E in Computer Science

Southwest University, Chongqing, China

2014 - 2018

B.E in Computer Science

RESEARCH EXPERIENCE

UNIVERSITY OF MARYLAND, COLLEGE PARK

College Park, MD

Research Assistant – Advisor: Wei Ai & Jing Liu

Sep. 2021 - present

- Discover latent and **distributional patterns** in text by combining LLMs with **causality-inspired** frameworks.
 - Investigate **bias** and **spurious correlation** mitigation for text **classification** and **LLM alignment**.
 - AI for Edu - LLM-based transcript analysis to identify effective teaching practices and measure teaching quality.
- Research Assistant** – Advisor: Louiqa Raschid & Vanessa Fries-Martinez Jan. 2023 - present
- Model correlations between geo-co-location and social media engagement using **statistical** & **LLM**-based methods.
 - Link online interaction patterns with offline community outcomes to promote pro-social behaviors.

JOHNS HOPKINS UNIVERSITY

Baltimore, MD

Full-time Research Developer – Advisor: Mark Dredze

Jun. 2020 - Aug. 2021

- Develop Twitter-based social mobility indices to study demographic differences in social distancing behavior. [[link](#)]
 - Apply embedding- and transformer-based models to discover e-cig brands and track sentiment shifts on social media.
- Research Assistant** – Mentor: Zach Wood-Doughty Nov. 2018 - May 2020
- Demographic inference on Twitter with distant supervision ([Package](#)); study linguistic variation across racial groups.

INDUSTRY EXPERIENCE

ADOBE RESEARCH

San Jose, CA

Research Scientist Intern – Mentor: Haoliang Wang

May 2025 - Aug. 2025

- Design style understanding and **personalized** suggestion generation through LVLM-based design pattern uncovering.
- **User-centric evaluation** for personalization using paired image comparisons to capture design preferences.

ADOBE RESEARCH

San Jose, CA

Research Scientist Intern – Mentor: Vishy Swaminathan, Gang Wu, & Xiang Chen

May 2024 - Aug. 2024

- **Skill development** for software scripting through code generation & offline simulations using **LLM-based agents**.

ADOBE RESEARCH

Remote

Research Scientist Intern – Mentor: Tong Yu & Haoliang Wang

May 2022 - Aug. 2022

- Learning knowledge from Critical Service Outage reports for auto-remediation.

- Adaptive Question Generation for Extractive Question Answering using In-context Learning.

SELECTED PUBLICATIONS

* denotes equal contribution

Paiheng Xu, Louiqa Raschid, Vanessa Frias-Martinez, Does Geo-co-location Matter? A Case Study of Public Health Conversations during COVID-19. *To appear in ICWSM 2026.*

Paiheng Xu, Gang Wu, Xiang Chen, Tong Yu, Chang Xiao, Franck Dernoncourt, Tianyi Zhou, Wei Ai, Vishy Swaminathan, Skill Discovery for Software Scripting Automation via Offline Simulations with LLMs. *arXiv:2504.20406*, 2025.

Paiheng Xu*, Xiaoyu Liu*, Junda Wu, Jiaxin Yuan, Yifan Yang, Yuhang Zhou, Fuxiao Liu, Tianrui Guan, Haoliang Wang, Tong Yu, Julian McAuley, Wei Ai, Furong Huang, Large Language Models and Causal Inference in Collaboration: A Comprehensive Survey. *Findings of NAACL 2025.*

Paiheng Xu*, Yuhang Zhou*, Bang An, Wei Ai, Furong Huang, GFairHint: Improving Individual Fairness for Graph Neural Networks via Fairness Hint. *Transactions on Knowledge Discovery from Data*, 2025

Paiheng Xu, Jing Liu, Nathan Jones, Julie Cohen, Wei Ai, The Promises and Pitfalls of Using Language Models to Measure Instruction Quality in Education. *NAACL 2024.*

Yuhang Zhou, **Paiheng Xu**, Xiaoyu Liu, Bang An, Wei Ai, Furong Huang, Explore Spurious Correlations at the Concept Level in Language Models for Text Classification. *ACL 2024.*

Paiheng Xu, David A. Broniatowski, Mark Dredze, Twitter Social Mobility Data Reveal Demographic Variations in Social Distancing Practices During the COVID-19 Pandemic. *Scientific Reports*, 2024.

Yuhang Zhou*, Jing Zhu*, **Paiheng Xu**, Xiaoyu Liu, Xiyao Wang, Danai Koutra, Wei Ai, Furong Huang, Multi-Stage Balanced Distillation: Addressing Long-Tail Challenges in Sequence-Level Knowledge Distillation. *Findings of EMNLP 2024.*

Paiheng Xu*, Zach Wood-Doughty*, Xiao Liu, Mark Dredze, Using Noisy Self-Reports to Predict Twitter User Demographics. *SocialNLP@NAACL*, 2021.

Paiheng Xu, Mark Dredze, David A. Broniatowski, The Twitter Social Mobility Index: Measuring Social Distancing Practices from Geolocated Tweets. *Journal of Medical Internet Research*, 2020.

Paiheng Xu, Likang Yin, Zhongtao Yue, Tao Zhou, On Predictability of Time Series. *Physica A: Statistical Mechanics, its Applications* 523, 345-351 (2019).

Paiheng Xu, Rong Zhang, Yong Deng, A Novel Visibility Graph Transformation of Time Series into Weighted Networks. *Chaos, Solitons & Fractals* 117, 201-208 (2018).

TECHNICAL SKILLS

| | |
|-----------|--|
| Languages | Python, MATLAB, R, C/C++, C#, L ^A T _E X, Markdown, Bash |
| Tools | Transformers, PyTorch, TRL, HF Agents, Scikit-learn, NLTK, Statsmodels, NetworkX |

HONOR & AWARDS

Dean's Fellowship, University of Maryland

National Scholarship, Ministry of Education, China

Outstanding Graduates, Southwest University

Meritorious Winner, Interdisciplinary Contest in Modeling (ICM)

TEACHING & SERVICE

Teaching Assistant: Introduction to Artificial Intelligence, Fall 2022

Program Committee: ICWSM 2022, 2025

Reviewer: ARR 2025; KDD 2025; TheWebConf 2025; TPAMI

Sub-Reviewer: ACL 2021; KDD 2022; SIGIR 2022-2025